

How to Deploy Ceph on OpenPOWER

Version 1.3

INTRODUCTION

This document along with referenced links describes a comprehensive set of instructions, rules, and automation tools for building an OpenPOWER-based platform for a Ceph cluster with Operational Management. This provides a standalone configuration for Ceph. For a Ceph cluster that is integrated with an OpenStack private compute cloud, see the private compute cloud documentation. The standalone Ceph configuration includes distinct nodes for the controller (where ceph-mon and opsmgr services reside) and ceph-osd.

The build process is broken out into a series of Steps listed below. Some of these steps are preparatory in nature. The final step is fully automated. At the completion of this step, a standalone Ceph cluster should be operational.

HIGH LEVEL DEPLOYMENT STEPS

EACH STEP BELOW IS DESCRIBED IN MORE DETAIL BELOW

1	Acquire the hardware
2	Choose deployment parameters
3	Prepare the deployment node
4	Rack the servers, and switches
5	Cable the systems
6	Deploy the cluster

STEP 1: ACQUIRE THE HARDWARE

Go [here](#) to view the Design Proposal for recipe required hardware.

<https://github.com/open-power-ref-design/standalone-ceph/blob/master/docs/design.pdf>

Go [here](#) for the Bill of Materials list of required parts:

<https://github.com/open-power-ref-design/standalone-ceph/blob/master/docs/bom.pdf>

If you do not already have the needed parts, please [contact](#) an IBM representative to assist you.

<https://www-01.ibm.com/marketing/iwm/dre/signup?source=MAIL-power&disable-Cookie=Yes>

STEP 2: CHOOSE YOUR DEPLOYMENT PARAMETERS

To facilitate faster automated configuration of the overall solution, collect together the following parameters before you start. This data will be edited into a configuration file which will in turn be used to automatically tune the infrastructure and add other needed software.

Parameter	Description	Example
Domain Name	The local domain where the cluster is being built.	mycompany.domain.com
Upstream DNS Servers	While a DNS server is configured within the cluster, upstream DNS servers need to be defined for names that cannot otherwise be resolved.	*4.4.4.4, 8.8.8.8 as default public upstream DNS servers
Deployment Node Host Name	What do you want to call your deployment node?	depnode
Management network IP address range	Management for the cluster happens on its own internal network. Labeled <i>ipaddr-mgmt-network</i> in the config.yml example below.	192.168.16.0/24

Data network IP address range	Private data network internal to the cluster. Labeled <i>ceph-public-storage</i> in the config.yml example below.	10.0.16.0/22
Management switch IP address	IP address of the management switch. Labeled <i>ipaddr-mgmt-switch</i> in config.yml in example below	192.168.16.20
Data switch IP address	IP address of the data switch on the rack. Labeled <i>ipaddr-data-switch</i> in config.yml in example below	192.168.16.25
External floating IP address	Floating ip-address that can be used to view the cluster GUI (Horizon) externally	10.0.16.50
Default login data	Both IDs and passwords	BMC network, OS Mgmt. network
Client-OS-Level	What operating system to install on the nodes	ubuntu-16.04.1-server-ppc64el

Go [here](#) to see more options in the config.yml file.

<https://github.com/open-power-ref-design/standalone-ceph/blob/master/config.yml>

STEP 3: PREPARE THE DEPLOYER NODE

The deployer node is used to obtain the latest software and deployment tools from GitHub and populate the cluster. The deployment node is not included in the Bill of Materials (BOM). You can establish the deployer node as a temporary or a permanent server. It can be any Power8-LC or x86 server with the following minimum characteristics:

- 2 cores and 32G RAM
- 1 Network Interface connection: 1G (Mgmt.).
- Ubuntu 16.04 LTS.

If you do not already have Ubuntu installed on the deployer node, you can obtain it from the following locations:

- Power8-LC servers: <https://www.ubuntu.com/download/server/power8>
- For x86 servers: <http://releases.ubuntu.com/16.04.1/>

STEP 4: RACK THE SERVERS, SWITCHES AND STORAGE

There are various compute, storage and networking components in an Ceph solution. Optimal server resources have been selected to attain the right balance between compute, storage and network I/O.

The two classes of servers in a standalone Ceph solution are:

Controller Nodes: These servers host the Ceph Monitor services and OpsMgr control services. The design is triple redundant and allows for scaling to larger (intra-rack) clusters. This solution has three [Stratton Power S821LC MTM 8001-12C](#) nodes serving this role. These nodes are referenced in the rest of the document with a prefix of 'Ctrl'.

Ceph OSD Nodes: These servers host the Ceph OSD (Object Storage Daemon) services. This solution has [Briggs Power S822LC MTM 8001-22C](#) nodes serving this role. These nodes are referenced in the rest of the document with a prefix of 'OSD'.

Racking the components

Place the intra-rack (leaf) network switches in **U24-U26** as follows:


- 10G data plane switches in **U25 and U26** (part 8831-S48)
- 1G management plane switch in **U24**. (part 7120-48E)

Place the OSD nodes at the bottom of the rack. Place the Ceph monitor nodes above the OSD nodes.

Ceph monitors should use different sets of PDUs for redundancy. OSD nodes should use different sets of PDUs for redundancy.

If more than 6 PDUs are needed, place 2 horizontal PDUs in 40U and 41U. Spine switches take priority over additional PDUs. If 40U and 41U are occupied by Spine Switches, place horizontal PDUs in next available slots.

RESULTING EXAMPLE: 12 NODE STANDALONE CEPH CLUSTER



STEP 5: CABLE THE SYSTEMS

Cabling of a Standalone Ceph solution consists of two distinct efforts. One is the network cabling and the other is powering the systems and switches. There are no external storage drawers so no separate storage cabling is required.

CABLE THE NETWORK


There are two networks in the Standalone Ceph solution. There is a management network and a data network each with its own dedicated switch. The ports marked IPMI and PXE ports go to the 1G management network/switch, and the ports marked Data go to the 10G data network/switch.

Take the network cables and connect the IPMI port on the server to the management switch at the location marked on the switch for that node. Perform the same step next with the port marked PXE. Finally take a third network cable and connect the port marked with Data on the server to the data switch at the location marked for that node.

Since the S822LC (8001-22C) "Briggs" and S821LC (8001-12C) "Stratton" server nodes are not identical we include below examples depicting ports for both models.


EXAMPLE: STRATTON NODE

Rear view of Stratton server with labels



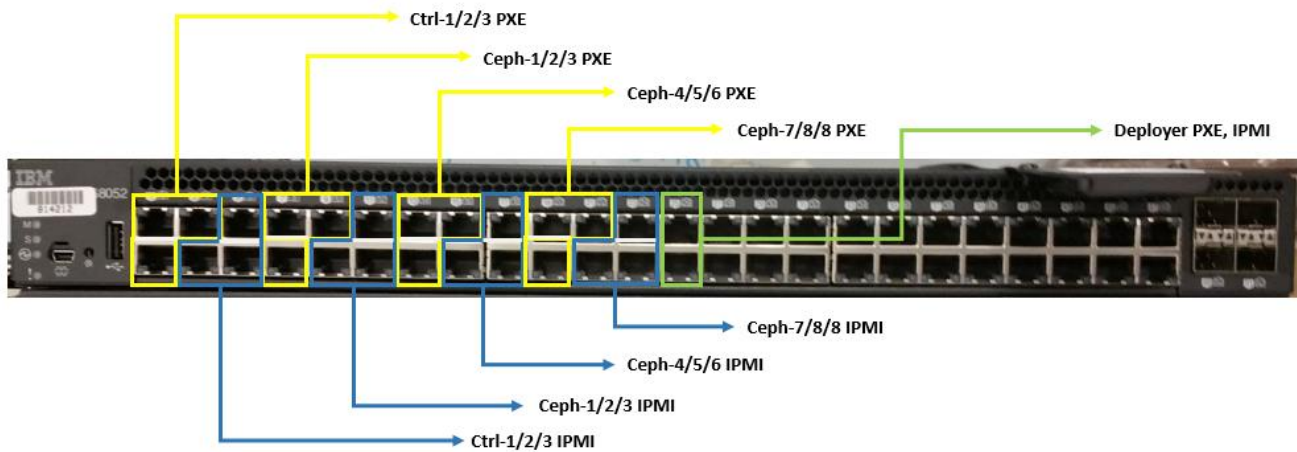
EXAMPLE: BRIGGS NODE

Rear view of Briggs server with labels




Once the server nodes are cabled please complete cabling connections on the Mgmt and Data switches.

Rear-view of management network switch.



Rear-view of Data Network Switch.



Repeat the above steps for the other Ceph-OSD and Ctrl nodes.

POWER CONSIDERATIONS

Careful consideration is required in cabling up the Power for the servers. Servers, storage drawers and switches should be distributed among different PDUs. The PDUs have to be provided independent Power drops to ensure that a single PDU failure or a single power line will not disrupt any of the services. For example say PDU1 and PDU3 are attached to the first

power line and PDU2 and PDU4 are attached to the second power line. To ensure that the Ceph services are always available, Ceph-1, Ceph-2 and Ceph-3 nodes should not all be connected to PDU1 and PDU3. Instead Ceph-1 could be connected to PDU1, Ceph-2 to PDU2 and Ceph-3 to PDU3. For servers with redundant power supplies, each supply should be connected to a separate path (PDU/line). The same goes for the redundant network switches. One should be connected to PDU1/PDU3 and the other to PDU2/PDU4. This will ensure the highest availability.

STEP 6: DEPLOY THE CLUSTER

This section covers the power on, initialization, configuration and installation of a Standalone Ceph cluster. This deployment kit provides an automated method to quickly and more predictably go from assembly to a tuned operational state of the cluster's infrastructure.

THE STEPS INVOLVED IN DEPLOYING A CLUSTER INCLUDE:

A	Obtain the default configuration file
B	Tailor the configuration file for your environment
C	Validate the configuration file
D	Provision the cluster
E	Configure OpenStack services.
F	Verify the deployment

The orchestration of the above steps is performed by a tool called cluster-genesis. Please refer to the user guide [here](http://cluster-genesis.readthedocs.io/en/latest/) for details on cluster-genesis.

<http://cluster-genesis.readthedocs.io/en/latest/>

OBTAIN THE DEFAULT CONFIGURATION FILE

The deployment automation (cluster-genesis) uses a config file to specify the target cluster configuration. The deployment tooling uses this YAML text file to specify the IP address locations of the managed switches and the system nodes attached to the switches as well as other useful details for the deployment process.

Go [here](#) for a copy of the Standalone Ceph configuration file.

<https://github.com/open-power-ref-design/standalone-ceph/blob/master/config.yml>

TAILOR THE CONFIGURATION FILE FOR YOUR ENVIRONMENT

The config.yml file contains a lot of configuration information. To enable a cluster tailored to your environment, edit the YAML file with the configuration parameters you collected in Step 2, replacing the **RED** text with your data. The image below zooms in on only those lines to edit.

Editable Portions of the config.yml file

```
~ ~ ~ ~ ~ bunch of licensing comment and YAML ~ ~ ~ ~ ~
ipaddr-mgmt-network: 192.168.16.0/24 ⚡ Type your management network range here.
ipaddr-mgmt-switch:
  rack1: 192.168.16.20 ⚡ Type your management switch IP address here.
ipaddr-data-switch:
  rack1:
    - 192.168.16.25 ⚡ Type your data switch1 IP address here.
    - 192.168.16.30 ⚡ Type your data switch2 IP address here.
external-floating-ipaddr: 10.0.16.50 ⚡ Type your external floating IP address here.

~ ~ ~ ~ ~ series of YAML and comments ~ ~ ~ ~ ~
networks:
  ceph-public-storage:
    description: Organization site or external network
```

```

addr: 10.0.16.0/22           ← Type your External Data IP address range here.
broadcast: 10.0.19.255      ← Type your External network broadcast address here.
gateway: 10.0.16.1         ← Type your external router address here.
dns-nameservers: 10.0.16.200 ← Type your DNS server IP address here.
dns-search: mycompany.domain.com ← Type your DNS domain name here
method: static
eth-port: osbond0

```

~ ~ ~ ~ ~ series of YAML and comments ~ ~ ~ ~ ~

node-templates:

controllers:

```

hostname: ctrl-             ← Type your controller-1 hostname here.
userid-ipmi: ADMIN         ← Type your IPMI user here.
password-ipmi: admin       ← Type your IPMI password here.
cobbler-profile: ubuntu-16.04.1-server-ppc64el ← Type your OS Level here

```

~ ~ ~ ~ ~ continuation of YAML and comments ~ ~ ~ ~ ~

VALIDATE THE CONFIGURATION FILE

To ensure that the format of the modified configuration file is valid, it is recommended that it gets validated by following these steps:

```

$ git clone git://github.com/open-power-ref-design/standalone-ceph
$ cd standalone-ceph
$ TAG=$(git describe --tags $(git rev-list --tags --max-count=1))
$ cd ..
$ apt-get install python-pip
$ pip install pyyaml
$ git clone git://github.com/open-power-ref-design-toolkit/os-services
$ cd os-services
$ git checkout $TAG
$ ./scripts/validate_config.py --file ../standalone-ceph/config.yml
$ cd ..

```

Before proceeding with the Provisioning step below, you will need to ensure that all the nodes in the cluster will have direct access to the internet. If the cluster is being configured in a private network without direct internet access, then it is recommended that the deployer node be provided internet access, and it acts as a NAT host to route all the nodes in the cluster.

PROVISION THE CLUSTER

Once the deployment cluster's configuration is finalized, the cluster can be provisioned by kicking off the automation tool (cluster-genesis). The initial steps are:

```
$ git clone git://github.com/open-power-ref-design-toolkit/cluster-genesis
$ cd cluster-genesis
$ TAG=$(git describe --tags $(git rev-list --tags --max-count=1))
$ git checkout $TAG
$ ./scripts/install.sh
$ source scripts/setup-env
$ gen deploy
```

If custom install images are required, more details of the steps to follow are found [here](#):

http://cluster-genesis.readthedocs.io/en/latest/OPCG_running_OPCG.html

CONFIGURE OPENSTACK SERVICES

The provisioning step above, runs for about 2 hours and completes the installation of the operating system on all nodes. Upon completion of the installation it then prepares the cluster (bootstrap) as well. The bootstrap step is automatically triggered when cluster-genesis is completed. At this point various OpenStack parameters need to be configured. To prepare for this phase, please collect the following information:

Parameter	Description	Example
Keystone Password	The password that will be used for the OpenStack authentication services	passw0rd
VRRP ID	Virtual Router ID that will be in the range of 1-255 and has to be unique across the network	202
USED IPs	The range of IP addresses in the private networks that are already taken, and cannot be assigned to	172.29.236.1,172.29.236.50 172.29.240.1,172.29.240.50 172.29.244.1,172.29.244.50

	nodes in the cluster. This includes the addresses assigned by cluster-genesis.	
--	--	--

For complete information on all the various options to configure the OpenStack deployment, please refer to the openstack-ansible documentation available [here](#).

<https://docs.openstack.org/project-deploy-guide/openstack-ansible/newton>

The minimal set of configuration options are provided below. To proceed with this, login to the first controller node (ctrl-1).

```
$ ssh ctrl-1
```

Edit the keystone stanza and set the desired keystone password, in the /etc/openstack_deploy/user_secrets.yml file.

```
$ vi /etc/openstack_deploy/user_secrets.yml
```

```
## Keystone Options
keystone_container_mysql_password:
keystone_auth_admin_token:
keystone_auth_admin_password: passw0rd
keystone_service_password:
keystone_rabbitmq_password:
```

The valid range for the next parameter (external virtual router ID) is 1-255 and it must be unique for each cluster. Edit the external virtual router ID, in the /etc/openstack_deploy/user_variables.yml file.

```
$ vi /etc/openstack_deploy/user_variables.yml
```

```
# Defines the default VRRP id used for keepalived with haproxy.
# Overwrite it to your value to make sure you don't overlap
# with existing VRRPs id on your network. Default is 10 for the
# external and 11 for the internal VRRPs
haproxy_keepalived_external_virtual_router_id: 202
# haproxy_keepalived_internal_virtual_router_id:
```

Next edit the /etc/openstack_deploy/openstack_user_config.yml file to reserve ip addresses generally for the 172 networks: .1-.50. This is done with the "used_ips" field described in /etc/openstack_deploy/openstack_user_config.yml.example. The following values can be placed just above the "global_overrides" field.

```
$ vi /etc/openstack_deploy/openstack_user_config.yml
```

```
used_ips:
- "172.29.236.1,172.29.236.50"
- "172.29.240.1,172.29.240.50"
- "172.29.244.1,172.29.244.50"
```

Now the cluster is ready to complete the final step of deployment:

```
$ cd os-services
$ ./scripts/create-cluster.sh 2>&1 | tee -a /root/create-cluster.out
```

Monitor the /root/create-cluster.out file for progress and indication of completion.

(OPTIONAL) CONFIGURE OPERATIONAL MANAGEMENT SERVICES

Part of the software stack deployed in the Controller nodes is a set of services that are collectively called "Operational Management" (or OpsMgr for short). This complements Ceph with popular DevOps tools that provide additional function to monitor availability and health of your cluster and to collect log information from all elements of the cluster and present key metrics that are useful for continued operations. These tools are, respectively, Nagios Core and the Elastic Stack. In addition, Operational Management also offers Hardware inventory and integration services as an extension to the OpenStack Horizon user interface. From that extension users can visualize all their hardware devices and launch to all integrated DevOps tools to perform further operational tasks.

For more information on Operational Management and to learn about optional configuration of its installed services and tools please consult the associated [README](#).

<https://github.com/open-power-ref-design-toolkit/opsmgr/blob/master/recipes/private-cloud-newton/README.rst>

VERIFY THE DEPLOYMENT

The create-cluster provisioning step above will run for several hours (~3). After its completion to verify that the Standalone Ceph cluster is operational, please follow the 'Verifying an install' section of the [README](#).

<https://github.com/open-power-ref-design/standalone-ceph/blob/master/README.rst>

REFERENCE LINKS

Go [here](#) for the Stratton Power S821LC MTM 8001-12C Redbook:

<http://www.redbooks.ibm.com/abstracts/redp5406.html?Open>

Go [here](#) for the Briggs Power S822LC MTM 8001-22C Redbook.

<http://www.redbooks.ibm.com/abstracts/redp5407.html?Open>

Go [here](#) for standalone-ceph repository:

<https://github.com/open-power-ref-design/standalone-ceph>

Go [here](#) for the OpenStack-Ansible documentation:

<https://docs.openstack.org/developer/openstack-ansible/newton>



© Copyright International Business Machines Corporation 2017

Printed in the United States of America June 2017

IBM, the IBM logo, and ibm.com are trademarks or registered trademarks of International Business Machines Corp., registered in many jurisdictions worldwide. Other product and service names might be trademarks of IBM or other companies. A current list of IBM trademarks is available on the Web at “Copyright and trademark information” at www.ibm.com/legal/copytrade.shtml.

NVLink is a trademark of the NVIDIA Corporation in the United States, other countries, or both.

The OpenPOWER word mark and the OpenPOWER Logo mark, and related marks, are trademarks and service marks licensed by OpenPOWER.

Other company, product, and service names may be trademarks or service marks of others.

All information contained in this document is subject to change without notice. The products described in this document are NOT intended for use in applications such as implantation, life support, or other hazardous uses where malfunction could result in death, bodily injury, or catastrophic property damage. The information contained in this document does not affect or change IBM product specifications or warranties. Nothing in this document shall operate as an express or implied indemnity under the intellectual property rights of IBM or third parties. All information contained in this document was obtained in specific environments, and is presented as an illustration. The results obtained in other operating environments may vary.

This document is intended for the development of technology products compatible with Power Architecture®. You may use this document, for any purpose (commercial or personal) and make modifications and distribute; however, modifications to this document may violate Power Architecture and should be carefully considered. Any distribution of this document or its derivative works shall include this Notice page including but not limited to the IBM warranty disclaimer and IBM liability limitation. No other licenses (including patent licenses), expressed or implied, by estoppel or otherwise, to any intellectual property rights are granted by this document.

THE INFORMATION CONTAINED IN THIS DOCUMENT IS PROVIDED ON AN “AS IS” BASIS. IBM makes no representations or warranties, either express or implied, including but not limited to, warranties of merchantability, fitness for a particular purpose, or non-infringement, or that any practice or implementation of the IBM documentation will not infringe any third party patents, copyrights, trade secrets, or other rights. In no event will IBM be liable for damages arising directly or indirectly from any use of the information contained in this document.

IBM Systems
294 Route 100, Building SOM4
Somers, NY 10589-3216

The IBM home page can be found at ibm.com®.

Version 1.3
12 June 2017